

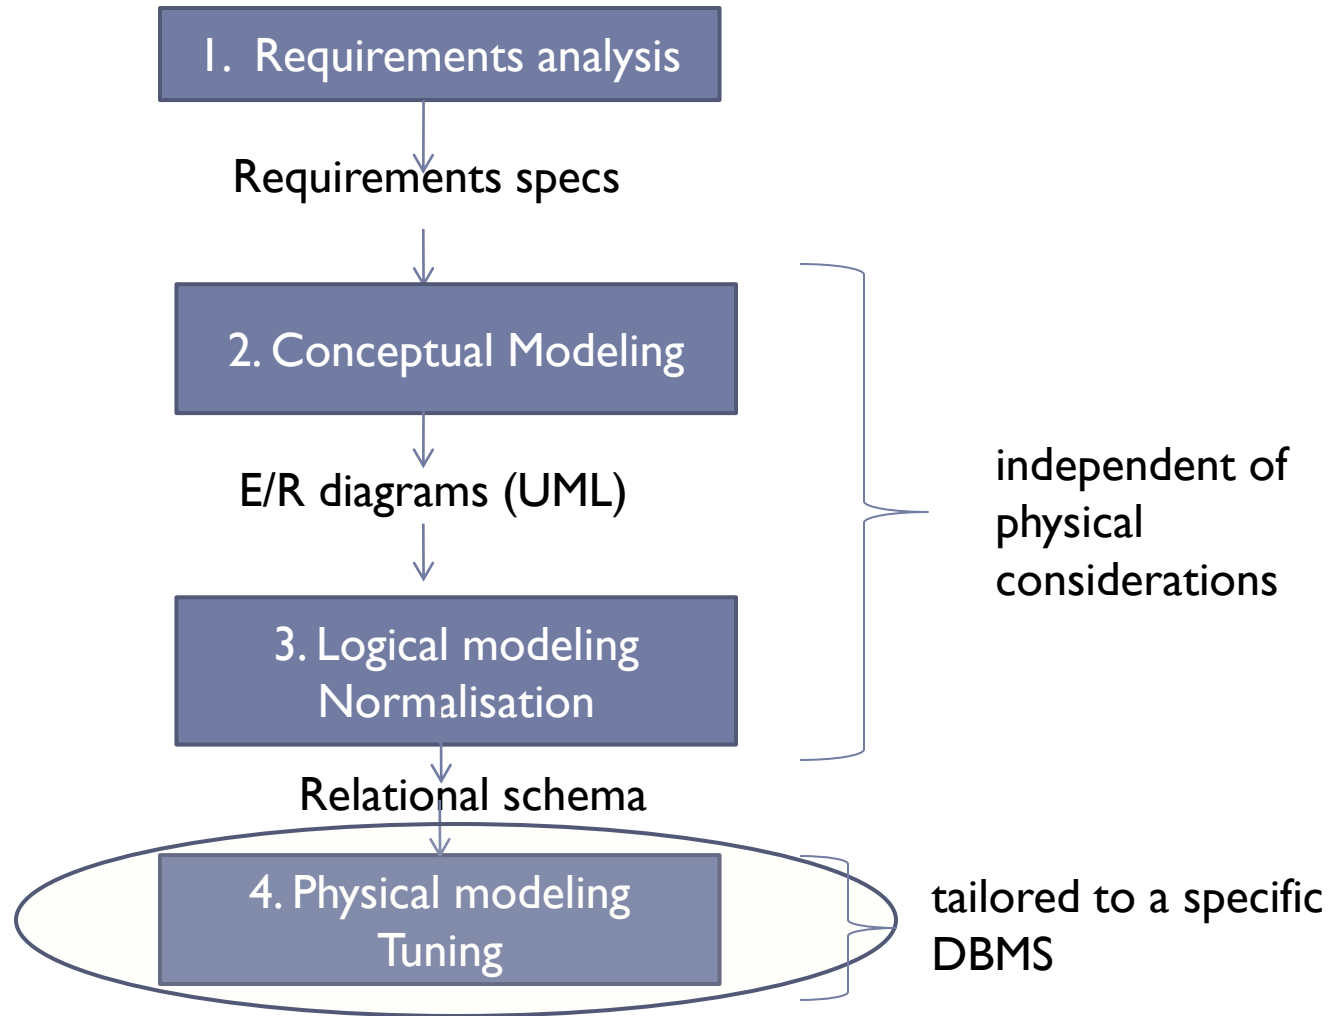


```
CREATE MATERIALIZED VIEW lecture8 AS  
SELECT lecture  
FROM Databases
```

```
WHERE topic = ' Physical Design ' ;  
INSERT INTO lecture8 VALUES (' Constraints, Triggers, Views');
```

Mihaela Elena Breabăn
© FII 2017-2018

Relational Database Design Methodology



Outline

- ▶ Declaring integrity constraints
- ▶ Triggers
- ▶ Views

Static integrity constraints

(1)

- ▶ **Restrict the possible states of the database**
 - ▶ Describe conditions that every instance of a database must satisfy
 - ▶ Avoid erroneous insertions, updates, deletions
 - ▶ Enforce data consistency
 - ▶ Tell the DBMS information that is useful for data storing and retrieval (query optimization)
- ▶ **Types**
 - ▶ Domain constraints
 - ▶ Non-null constraints
 - ▶ Keys – unique constraints
 - ▶ Referential integrity
 - ▶ General constraints at tuple level
 - ▶ Assertions

Static integrity constraints (2)

- ▶ **Declaration**
 - ▶ At (table) creation(CREATE TABLE)
 - ▶ After (table) creation (ALTER TABLE)
- ▶ **Validation**
 - ▶ At every DML statement
 - ▶ At the end of the transaction

Integrity constraints over 1 attribute

Inline declaration

CREATE TABLE *table_name* (

a1 type **not null**, -- does not allow null entries

a2 type **unique**, --candidate key consisting of one attribute

a3 type **primary key**, -- primary key consisting of one attribute, implies {not null, unique}

a4 type **references** *table_name2* (*b1*), --foreign key consisting of one attribute

a5 type **check** (*condition*) – the condition is a Boolean expression built with attribute *a5*: (*a5*<11 and *a5*>4), (*a5* between 5 and 10), (*a5* in (5,6,7,8,9,10))...
)

Integrity constraints over several attributes

Out-of-line declaration

CREATE TABLE *table_name* (

a1 type,

a2 type,

a3 type,

a4 type,

primary key (*a1,a2*), -- primary key consisting of two (or more)attributes

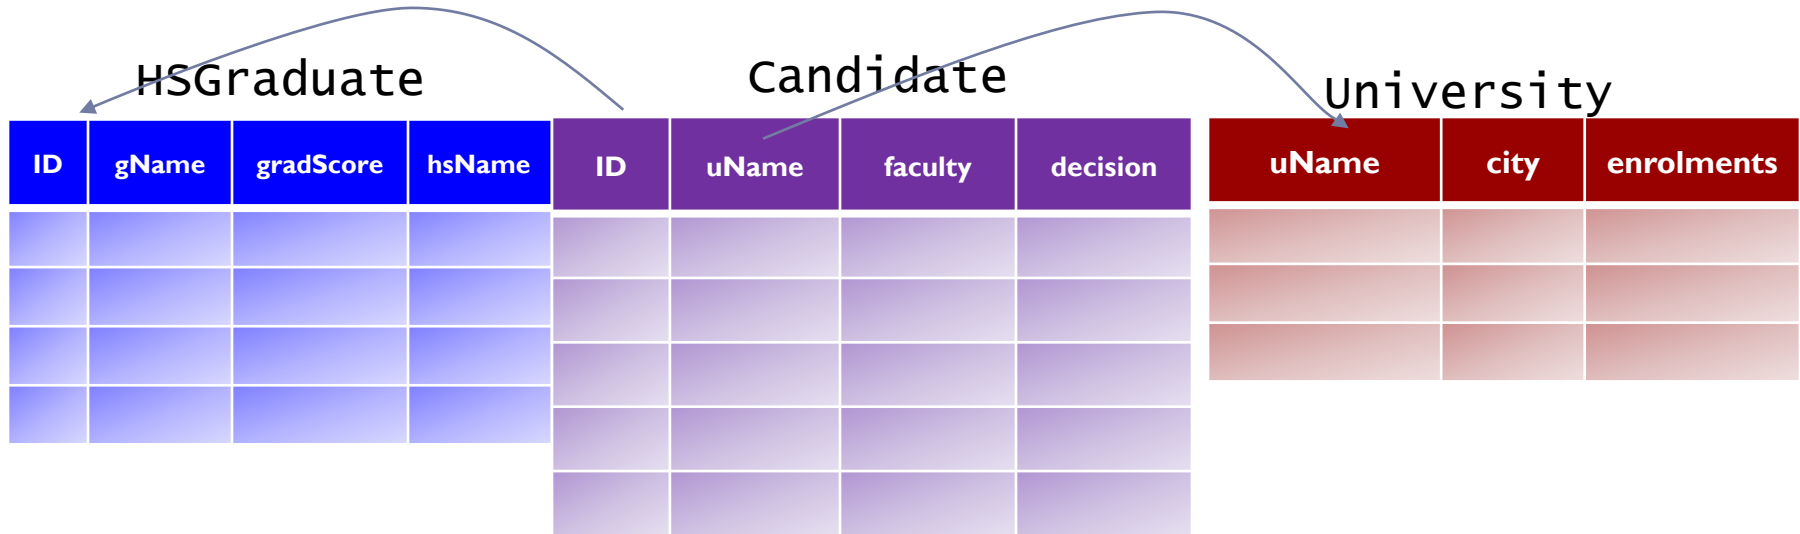
unique(*a2,a3*), -- candidate key consisting of two (or more)attributes

check (*condiție*), -- a Boolean expression built over several attributes:
($(a1+a3)/2 \geq 5$)

foreign key (*a3,a4*) **references** *table_name2*(*b1,b2*) – multi-valued foreign key
)

Referential integrity

Definitions



- ▶ *Referential integrity from R.A to S.B:*
 - ▶ Every value in column A of table R must appear in column B of table S
 - ▶ A is called a foreign key
 - ▶ B must be declared as primary key or unique for table S
- ▶ There may exist multi-valued foreign keys

Referential integrity Validation

- ▶ Statements that may generate violations:
 - ▶ Insertions in R
 - ▶ Deletions in S
 - ▶ Updates of R.A or S.B
- ▶ Special actions that can be enforced:
 - ▶ At deletions in S:
ON DELETE RESTRICT (by default) | **SET NULL** | **CASCADE**
 - ▶ At updates on S.B:
ON UPDATE RESTRICT (by default) | **SET NULL** | **CASCADE**

Referential integrity

egg or chicken?

```
CREATE TABLE chicken (cID INT PRIMARY KEY,  
                        eID INT REFERENCES egg(eID));  
CREATE TABLE egg(eID INT PRIMARY KEY,  
                  cID INT REFERENCES chicken(cID));
```

Referential integrity

egg or chicken?

```
CREATE TABLE chicken (cID INT PRIMARY KEY,  
                        eID INT REFERENCES egg(eID));  
CREATE TABLE egg(eID INT PRIMARY KEY,  
                  cID INT REFERENCES chicken(cID));
```

```
CREATE TABLE chicken(cID INT PRIMARY KEY, eID INT);  
CREATE TABLE egg(eID INT PRIMARY KEY, cID INT);
```

```
ALTER TABLE chicken ADD CONSTRAINT chickenREFegg  
  FOREIGN KEY (eID) REFERENCES egg(eID)  
  DEFERRABLE INITIALLY DEFERRED; -- Oracle
```

```
ALTER TABLE egg ADD CONSTRAINT eggREFchicken  
  FOREIGN KEY (cID) REFERENCES chicken(cID)  
  DEFERRABLE INITIALLY DEFERRED; -- Oracle
```

```
INSERT INTO chicken VALUES(1, 2);  
INSERT INTO egg VALUES(2, 1);  
COMMIT;
```

How do you solve insertions if the constraints are validated at each statement?

What about table drops?

Assertions

-defined in the SQL standard -

```
create assertion Key
```

```
check ((select count(distinct A) from T) =  
       (select count(*) from T));
```

```
create assertion ReferentialIntegrity
```

```
check (not exists (select * from Candidate  
                  where ID not in (select ID from HSGraduate)));
```

Integrity constraints

DBMS implementations

- ▶ Postgres, SQLite, Oracle, MySQL(innodb) implement and validate all constraints above
- ▶ No DBMS allows queries in the check constraint (deviation from the SQL standard)
- ▶ No DBMS implements assertions – their functionality can be provided by triggers

...DEMO...
(file *constraints.sql*)

Triggers

Dynamic constraints

- ▶ Monitor all the changes in a database, check conditions and initiate actions
- ▶ *Event-condition-action* rules
 - ▶ Bring within the DBMS elements from the application logic
 - ▶ Enforce constraints that cannot be expressed otherwise
 - ▶ Are expressive
 - ▶ May implement repairing actions
- ▶ The implementation may differ among DBMSs, the examples in the presentation are in accordance with the SQL standard

Triggers

Implementation

Create Trigger *name*

Before|After|Instead Of *event*

[*Referenced-variables*]

[**For Each Row**] -- the action is executed for each row altered by the event (row vs. statement types)

[**When (condition)**] – a boolean expression – exactly as within the WHERE clause in SQL queries

action -- în standardul SQL e o comandă SQL, în SGBD-uri poate fi bloc procedural

▶ *event:*

- ▶ **INSERT ON** *table*
- ▶ **DELETE ON** *table*
- ▶ **UPDATE [OF *a1,a2,...*] ON** *table*

▶ *Referenced-variables* (they are declared and then used in *condition* and *action*):

- ▶ **OLD TABLE AS** *var*
- ▶ **NEW TABLE AS** *var*
- ▶ **OLD ROW AS** *var* – only for *DELETE, UPDATE*
- ▶ **NEW ROW AS** *var* – only for *INSERT, UPDATE*

} Only for ROW triggers

Triggers

Example

- ▶ Referential integrity from R.A to S.B implementing cascade deletes

```
Create Trigger Cascade_Deletes
After Delete On S
Referencing Old Row As O
For Each Row
[ no conditions ]
Delete From R Where A = O.B
```

```
Create Trigger Cascade_Deletes
After Delete On S
Referencing Old Table As OT
[ For Each Row ]
[ no condition ]
Delete From R Where
    A in (select B from OT)
```

Triggers

Problems to consider

- ▶ Several triggers are activated simultaneously: which one is the first to be executed?
- ▶ The trigger action activates other triggers: chaining that can create cycles

Triggers Implementation

▶ Postgres

- ▶ The closest to the standard
- ▶ implements row+statement -> {old,new}x{row,table} variables
- ▶ The syntax suffers some changes

▶ SQLite

- ▶ Only row (no old/new table)
- ▶ Are executed after each row modification

▶ MySQL


- ▶ Only row (no old/new table)
- ▶ Are executed after each row modification
- ▶ Allows only one trigger per event per table

▶ Oracle

- ▶ Implements the standard: row+statement, some syntax changes
- ▶ Instead-of triggers are allowed only for views
- ▶ Allows the use of procedural blocks
- ▶ Introduce restrictions to avoid cycling

...DEMO...
(file *triggers.sql*)

Views



Motivation

- ▶ Controlled access to the data:
 - ▶ hide data from specific users
 - ▶ restrict DML statements
- ▶ Reduce the complexity of queries
- ▶ Real applications tend to use many views

Definition and use

- ▶ A view is actually a stored query formulated over tables or other views
- ▶ Its schema is generated based on the schema of the query result
- ▶ Conceptually, a view is queried just like a table
- ▶ In reality, a query over a view is rewritten by replacing the name of the view with the query defining the view; query optimization takes place, as implemented by the DBMS
- ▶ Syntax

Create View *view_name* [(a1,a2,...)] **As** <select_statement>

Modifying views

- ▶ Although the main operation/statement executed on a view is querying it, they may allow DML operations
- ▶ Modification commands on views must be rewritten as modification commands over base tables
 - ▶ Usually is simple
 - ▶ Sometimes several ways exist
- ▶ Example 1
 - ▶ $R(A,B), V(A)=R[A];$
 - ▶ `INSERT INTO V VALUES(3);`
- ▶ Example 2
 - ▶ $R(N), V(A)=avg(N),$
 - ▶ `UPDATE V SET A=7;`

Modifying views

Approaches

1. The view owner must rewrite all DML statements launched on the view as DML statements on the tables using the **INSTEAD OF** trigger
 - ▶ Covers all the cases
 - ▶ Guarantees data consistency
2. The SQL standard defines the existence of (inherently) **updatable views**:
 - ▶ The view must be created based on a single table T
 - ▶ The attributes in T that are not used in the view may have NULL entries or DEFAULT values
 - ▶ There is no aggregation used: no GROUP BY, no DISTINCT keyword

Materialized views

Create Materialized View *V* [*a1,a2,...*] **As** <select_statement>

- ▶ A new table *V* is created with schema defined by the select statement
- ▶ The tuples result of the query are inserted into *V*
- ▶ Any query on *V* is executed directly on the *V* table

- ▶ Advantages:
 - ▶ Specific to regular views + increased query execution speed
- ▶ Disadvantages:
 - ▶ *V* may increase in size
 - ▶ Any DMLs on the base tables require modifications on *V*
 - ▶ Any DMLs on *V* still must be translated into DMLs on base tables

How decide if materialize?

- ▶ Data size
- ▶ Query complexity
- ▶ The number of queries on the view
- ▶ The number of modifications on the base tables and the possibility of incrementally updating the view
- ▶ Trade-off the time for query processing over the time for propagating updates on the view when DML commands occur the base tables

...DEMO...
(file *views.sql*)

Bibliografie

- ▶ Hector Garcia-Molina, Jeff Ullman, **Jennifer Widom**: *Database Systems: The Complete Book (2nd edition)*, Prentice Hall; (June 15, 2008)
- ▶ Oracle:
 - ▶ http://docs.oracle.com/cd/B28359_01/server.111/b28310/general005.htm
 - ▶ <http://www.oracle-base.com/articles/9i/MutatingTableExceptions.php>
 - ▶ http://www.dba-oracle.com/t_avoiding_mutating_table_error.htm